

How the world's collective attention is being paid to a pandemic: COVID-19 related 1-gram time series for 24 languages on Twitter

Thayer Alshaabi,^{1,*} Michael V. Arnold,^{1,*} Joshua R. Minot,^{1,*} Jane Lydia Adams,¹ David Rushing Dewhurst,¹ Andrew J. Reagan,^{2,1} Roby Muhamad,³ Christopher M. Danforth,^{1,4} and Peter Sheridan Dodds^{1,4,†}

¹*Computational Story Lab, Vermont Complex Systems Center,
MassMutual Center of Excellence for Complex Systems and Data Science,
Vermont Advanced Computing Core, University of Vermont, Burlington, VT 05401.*

²*MassMutual Data Science, Amherst, MA 01002.*

³*International Class Program, Faculty of Social and Political Sciences, University of Indonesia, Jakarta, Indonesia.*

⁴*Department of Mathematics & Statistics, University of Vermont, Burlington, VT 05401.*

(Dated: March 31, 2020)

In confronting the global spread of the coronavirus disease COVID-19 pandemic we must have coordinated medical, operational, and political responses. In all efforts, data is crucial. Fundamentally, and in the possible absence of a vaccine for 12 to 18 months, we need universal, well-documented testing for both the presence of the disease as well as confirmed recovery through serological tests for antibodies, and we need to track major socioeconomic indices. But we also need auxiliary data of all kinds, including data related to how populations are talking about the unfolding pandemic through news and stories. To in part help on the social media side, we curate a set of 1000 day-scale time series of 1-grams across 24 languages on Twitter that are most ‘important’ for March 2020 with respect to March 2019. We determine importance through our allotaxonomic instrument, rank-turbulence divergence. We make some basic observations about some of the time series, including a comparison to numbers of confirmed deaths due to COVID-19 over time. We broadly observe across all languages a peak for the language-specific word for ‘virus’ in January followed by a decline through February and a recent surge through March. **The world’s collective attention dropped away while the virus spread out from China.** We host the time series on Gitlab, updating them on a daily basis while relevant. Our main intent is for other researchers to use these time series to enhance whatever analyses that may be of use during the pandemic as well as for retrospective investigations.

I. OVERVIEW

Understanding how major disasters affect the wellbeing of populations both in real time and historically is of paramount importance. We especially need real-time measurement to enable policy makers in health systems and government to gauge the immediate situation and evaluate scenarios, and for researchers to model possible future trajectories of social systems.

In this short piece, we describe how we select languages and 1-grams relevant to the time period of the present COVID-19 pandemic; show example time series plots for ‘virus’, including a visual comparison with COVID-19 confirmed case and death numbers; and describe the data sets, figures, and visualizations for 24 languages that we share online.

Rank	Language	Code	Rank	Language	Code
1	English	en	13	Hindi	hi
2	Spanish	es	14	Persian	fa
3	Portuguese	pt	15	Urdu	ur
4	Arabic	ar	16	Polish	pl
5	Korean	ko	17	Catalan	ca
6	French	fr	18	Dutch	nl
7	Indonesian	id	19	Tamil	ta
8	Turkish	tr	20	Greek	el
9	German	de	21	Swedish	sv
10	Italian	it	22	Serbian	sr
11	Russian	ru	23	Finnish	fi
12	Tagalog	tl	24	Ukrainian	uk

TABLE I. The 24 languages for which we provide COVID-19 related Twitter time series. See Tab. A1 for total counts of 1-grams per language in our data set from 2019/09/01 through to 2020/03/23.

II. SELECTION OF LANGUAGES AND 1-GRAMS

Our primary aim here is to generate a particular data stream that may be of help to other researchers: A principled set of 1-gram time series across major languages used on Twitter and news-relevant for March,

* The first three authors contributed as a team and are listed in alphabetical order.

† Corresponding author: peter.dodds@uvm.edu

2020. Our work is complementary to extant efforts to enable research on the COVID-19 pandemic [1, 2] by gathering and sharing epidemiological data [3, 4], economic data, and internet and social media data [5–7].

We base our curation on our work in two of our previous papers [8, 9], and we draw from a database of approximately 10% of all tweets from 2008/09/09 to present.

Our process of obtaining salient 1-grams for March 2020 comprises two steps.

First, in [8], we used the Language Identification (LID) model FastText [10, 11] to evaluate all tweets in our historical archive, finding over 150 languages. We subsequently extracted GMT day-scale Zipf distributions for 1-, 2-, and 3-grams along with day-scale n -gram time series. We focus here on 1-grams. We preserve case where applicable, do not apply any stemming, and include hashtags, handles, and emojis.

Besides analyzing all tweets (AT), we also separately process what we call organic tweets (OT): All Twitter messages which are original. Organic tweets exclude retweets while including all added text for quote tweets. In doing so, we are able to carry through a measure of spreadability for all n -grams. The key threshold we use for spreading is the naive one from biological and social contagion models: When a 1-gram appears in more retweeted than organic material, we view it as being socially amplified.

Here, we take 24 of the most commonly used languages on Twitter in 2019, with the provision that we are able to parse them into 1-grams. For the time being, we are unable to reliably parse Japanese, Thai, and Chinese, the 2nd, 6th, and 13th most common languages. We exclude all tweets not assigned a language with sufficient confidence (an effective 4th ranked collection) and choose to include Ukrainian (29th) over Cebuano (28th). We list the 24 languages by overall usage frequency in Tab. I.

Second, we compare usage of 1-grams in March of 2020 with March 2019 to determine which 1-grams have become most elevated in relative usage. We do so by using rank-turbulence divergence [9], an instrument for comparing any pair of heavy-tailed size distributions of categorical data. Other well-considered divergences will produce similar lists.

For each language, we take Zipf distributions for each of the first 21 days of March 2020, and compare them with the Zipf distributions of 52 weeks earlier. We use rank-turbulence divergence with the parameter α set to 1/3 as this provides a reasonable fit to the lexical turbulence we observe [9, 12].

For an example, we show in Fig. 1 an allotaxonograph for Italian comparing 2019/03/19 and 2020/03/19. For ease of plotting, we have further chosen to compare the subset of words containing latin characters only. Words associated with the pandemic dominate the contributions from 2020/03/19. On the right side of the allotaxonograph, we see ‘Coronavirus’, ‘virus’, ‘Bergamo’, ‘pandemia’, and ‘morti’.


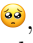
We next combine rank-turbulence divergence contri-

butions for all 1-grams across the first 21 days of March, and rank 1-grams in descending order.

Finally, for each language, and for each of the top 1-grams we have identified, we extract three day-scale time series starting on 2019/09/01: daily (GMT) counts, ranks, and normalized frequencies.

We repeat all of the above steps for 1-grams derived from organic tweets (OT).

We show the resulting top 20 March-2020-specific 1-grams for the 24 languages in Tabs. II and III. Overall, we see that the lists are dominated by language specific words for coronavirus virus, quarantine, pandemic, testing, and spreading. Not every language is pandemic focused however with Hindi, Tamil, and Finnish being three examples.

The microbe emoji, , makes the top 20 in only Spanish and German. By contrast, and according to the measurements we have used here, the worried face emoji, , has become important across many languages in March 2020 relative to March 2019. It would be natural to see this emoji as being pandemic-related but in fact, we see from time series that the worried emoji has slowly been increasing in usage over time for several years (determining the reasons for which we will leave for a separate line of inquiry).

We emphasize that with our approach, we do not explicitly determine whether or not a 1-gram is relevant to COVID-19. While the pandemic is the global story of March 2020, there have of course been other major events and moments in popular culture around the world. For example, for the United States, the ongoing democratic primary leads to the 1-gram in English Twitter of ‘Biden’ being prominent.

We also observe variations in punctuation and grammatical structures. All of these non-pandemic-related elements may of course be filtered out for individual languages by hand. While we do not have the capability and bandwidth across our team to perform such specific tasks, the degree to which the pandemic is being discussed on Twitter is of great interest in itself, and our data set will allow for such examination.

Further, most languages have a strong degree of geographic specificity (e.g., Finnish for Finland, Portuguese for Brazil), and we have not filtered for precise geolocation. English, Spanish, Arabic, and French are some of the more geographically distributed languages.

III. EXAMPLE TIME SERIES

We provide and briefly consider two sets of sample time series based on our data set.

Across Figs. 2 and 3, we plot usage rank time series for the word ‘virus’ translated as appropriate in to each of the 24 languages. (See Figs. A1 and A2 for normalized rates of usage versions.)

For each language, we show the daily (Zipfian) rank for ‘virus’ in the main panel of each plot. The pale disk




FIG. 1. Allotaxonograph using rank-turbulence divergence for Italian word usage on March 19, 2019 versus March 19, 2020. For this visualization, we consider the subset of 1-grams that are formed from latin characters. The right hand sides of the rank-rank histogram and the rank-turbulence contribution list are dominated by COVID-19 related terms. See [9] for a full explanation of our allotaxonometric instrument.

highlights the date of maximum observed rate. In the secondary time series at the top of each panel, we show the relative fraction of 1-gram contained in all tweets (RT) versus organic tweets (OT). When the RT/OT balance exceeds 50%, we shade the background to indicate that the 1-gram is being spread (e.g., retweeted) more than freshly tweeted.

Alone, the highest ranks for ‘virus’ show the enormity of the pandemic. While a common enough word in normal times, ‘virus’ has reached into the top 100 ranks across many languages, a region that we have elsewhere referred to as the realm of lexical ultrafame [13]. Normally only the most basic function words of a language will populate the top 100 ranks.

In the last few months, we have seen ‘virus’ rise as high as $r=24$ in Indonesian (2020/01/26), $r=27$ in Polish (2020/03/11), $r=29$ in Urdu (2020/03/22), $r=44$ in German (2020/03/14), and $r=83$ in English (2020/03/13).

In terms of the shapes of the time series for ‘virus’, most languages show a late January peak consistent with the news from China of a novel coronavirus disease

spreading in Wuhan. The subsequent drop in usage rate across most of the 24 languages reflects a global decline in attention being paid to the outbreak.

The Italian time series for ‘virus’ in Fig. 2J shows an abrupt jump about three quarters of the way through February, strikingly just after a drop in RT/OT balance. Persian has a similar shock jump just after midway of February (Fig. 2B).

A region currently experiencing a major outbreak is Catalonia. We see in Fig 3E that ‘virus’ in Catalan shows no early January peak like most of the other 23 languages, suggesting that even the initial news from China did not have great impact.

One of the major problems we face with the COVID-19 pandemic is the unevenness of testing across the world. South Korea and Iceland have tested early and extensively while the United States’s testing has been uncoordinated and slow to expand.

Urdu’s heightened time series for ‘virus’ (Fig 3C) would seem especially concerning given low numbers coming out of Pakistan which, as of 2020/03/24, had reported 1,063 cases and 8 deaths [3]. For Indonesia,

English	Spanish	Portuguese	Arabic	Korean	French
coronavirus	cuarentena	corona	كود	／	:
virus	virus	Daniel	خضم	／	virus
pandemic	pandemia	#BBB20	كورونا	코로나	quarantaine
quarantine	casos	quarentena	👉	한승우	2020
Biden	2020	babu	إستخدم	#영웅	confinement
corona	contagio	:	تخفيض	마스크	masques
Corona	infectados	virus	نمشی	승우	l'épidémie
outbreak	Italia	Thelma	قسیمه	2020	🤔
tested	medidas	Marcela	#كورونا	#VICTON	l'épidémie
#BestMusicVideo	maskarillas	Babu	فیروس	#빅톤	propagation
cases	contagios	🤔	کویون	김태형	#Municipales2020
Virus	🦠	daniel	قوغا	영웅	pa
@JoeBiden	contagiados	prior	سنتریویت	G	pandémie
2020	síntomas	Manu	کلوسیت	하울링	gel
🤔	epidemia	Ivy	اند	2020년	grippe
testing	propagación	paredão	نون	서비스	sanitaire
Wuhan	@alferdez	Rafa	الوباء	정국	Corona
#Master	🤔	marcela	نسناس	127	cas
@sidharth_shukla	salud	Corona	الكود	달고나	mesures
sanitizer	OMS	virus	ماکس	신천지	épidémie
Indonesian	Turkish	German	Italian	Russian	Tagalog
virus	@RTErdogan	.	virus	🤔	na
🤔	#ÇağlarErtuğrul	Virus	quarantena	карантин	🤔
2020	virüs	2020	contagio	Конституцию	virus
🤔	Ayşe	#Griechenland	intensiva	🤔	2020
Virus	virüsü	,	medici	2020	tiktok
masker	2020	🤔	,	поправки	🤔
pasien	Kerem	Grenze	2020	вирус	🤔
wabah	vaka	Klopapier	terapia	🤔	health
positif	@drfahrettinkoca	Quarantäne	ospedali	вируса	Heart
penyebaran	#NeslihanAtagül	Ausbreitung	🤔	нефть	Ken
ak	@avabdullahguler	Griechenland	misure	поправок	workers
Pasien	@yilmaztunc	infiziert	emergenza	Конституции	🤔
hand	Virüs	griechischen	Lombardia	эпидемии	Nak
China	@abdulhamitgul	🤔	.	заражения	gobierno
@collegemenfess	Sağlık	Grenzen	infermieri	конституцию	🤔
kesehatan	@saglikbakanligi	Krise	mascherine	маски	@BobOngHugots
penularan	Corona	👉	sanitaria	масок	Alab
ping	salgını	🦠	mascherina	Италии	DOH
pencegahan	karantinaya	Italien	corona	👉	ol
rumah	corona	👉	sanità	заболевших	mask

TABLE II. Top 20 (of 1,000) 1-grams for our top 12 languages for the first three weeks of March 2020 relative to a year earlier. Our intent is to capture 1-grams that are topically and culturally important during the COVID-19 pandemic. While overall, we see pandemic-related words dominate the lists across languages, we also find considerable specific variation. Words for virus, quarantine, protective equipment, and testing show different orderings (note that we do not employ stemming). Unrelated 1-grams but important to the time of March 2020 are in evidence; the balance of these are important for our understanding of how much the pandemic is being talked about. Some 1-grams such as punctuation represent functional changes in the use of Twitter across languages; we include them nonetheless. To generate these lists we use the allotaxonomic method of rank-turbulence divergence to find the most distinguishing 1-grams (see Sec. II, Fig. 1, and Ref. [9]). We note that these tables are images and cannot be copied; see the Sec. IV for data download sites.

Hindi	Persian	Urdu	Polish	Catalan	Dutch
यरस	ویروس	وائرس	@AndrzejDuda	virus	Bomboclaat
@SaintRampalJiM	قرنطینه	خلیل	😞	.	:
Asharamji	ماسک	سرمد	testów	confinament	bomboclaat
2020	@weareoneEXO	ماروی	PAD	mesures	virus
@MimrotReena	بهداشت	مرضی	wirus	55555555...	RT
FB	@B_hundred_Hyun	قمر	@M_K_Blonska	sanitari	@rivm
Rampal	شیوع	جسم	2020	crisi	besmet
Maharaj	@layzhang	عورت	Duda	sanitaris	tu
@Dr_Uditraj	میتلا	الرحمان	zdrowia	😞	😞
Saint	بیمارستان	ہے	wirusa	2020	RIVM
@Oplndia_in	قم	چینی	@SzumowskiLukasz	Perpinyà	2020
God	بیماری	خاقان	Dudy	π	.
@ravikishann	چین	ماسک	Włoch	casos	Griekse
@gmner_gkp	ایتلا	تدابیر	Włoszech	tancar	Italië
channel	#EXO	ایران	chorych	Itàlia	verspreiding
@HemantSorenJMM	الحمدلله	الرحمن	przypadków	salut	getest
♦	بیماران	2020	🤔	Temperatura	maatregelen
Follow	پزشکی	جیو	epidemii	ໄກ໊	testen
Sadhna	😞	شاید	testy	simptomes	patiënten
@beingarun28	کادر	🇮🇳	szpitala	@salutcat	
Tamil	Greek	Swedish	Serbian	Finnish	Ukrainian
@actorvijay	.	Italien	.	amg	🐱
ரஜ	σύννορα	😞	virus	😞	🏆
@Jagadishbliss	'Εβρο	@AgnesWold	😞	manaa	😞
👉	2020	2020	#Srbija	via	BARK
🔥	κρούσματα	gränsen	😞	🌐	ak
😞	μέτρα	Ak	mere	ak	2020
@rajinikanth	*	Turkiet	Ja	@MarinSanna	Зеленський
@Samaniyantweet	Italia	ak	📌	🕊	@censor_net
📌	Ερντογάν	virus	@anabrnbic	fodase	!!
@anirudhofficial	εκκλησίες	Grekland	2020	3 7	👉
😞	υγείας	stänger	maske	.	@sashkof2
@mohandreamer	Τουρκία	@BengtHojer	amei	@THLorg	@EspressoTV
தளபத	μετανάστες	@hanifbali	da	tu	Italiï
ளபத	μάσκες	stänga	👉	:	⚡
✨	κλείσουν	Kina	#SNS	yhh	—
👉	πρόσφυγες	@Folkhalsomynd	Æ	yh	Зеленського
👉	μάσκα	repp	Italiji	@MariaOhisalo	😞
#ThalaAjith	τεστ	kris	virusa	🌿	👉
😞	συνόρων	sjukvården	Брнабић	👉	уряд
😞	@PrimeministerGR	spridning	link	sim	ОРДЛО

TABLE III. Continuing on from Fig. II: Top 20 1-grams for the second 12 of 24 languages we study for March 2020 relative to March 2019.

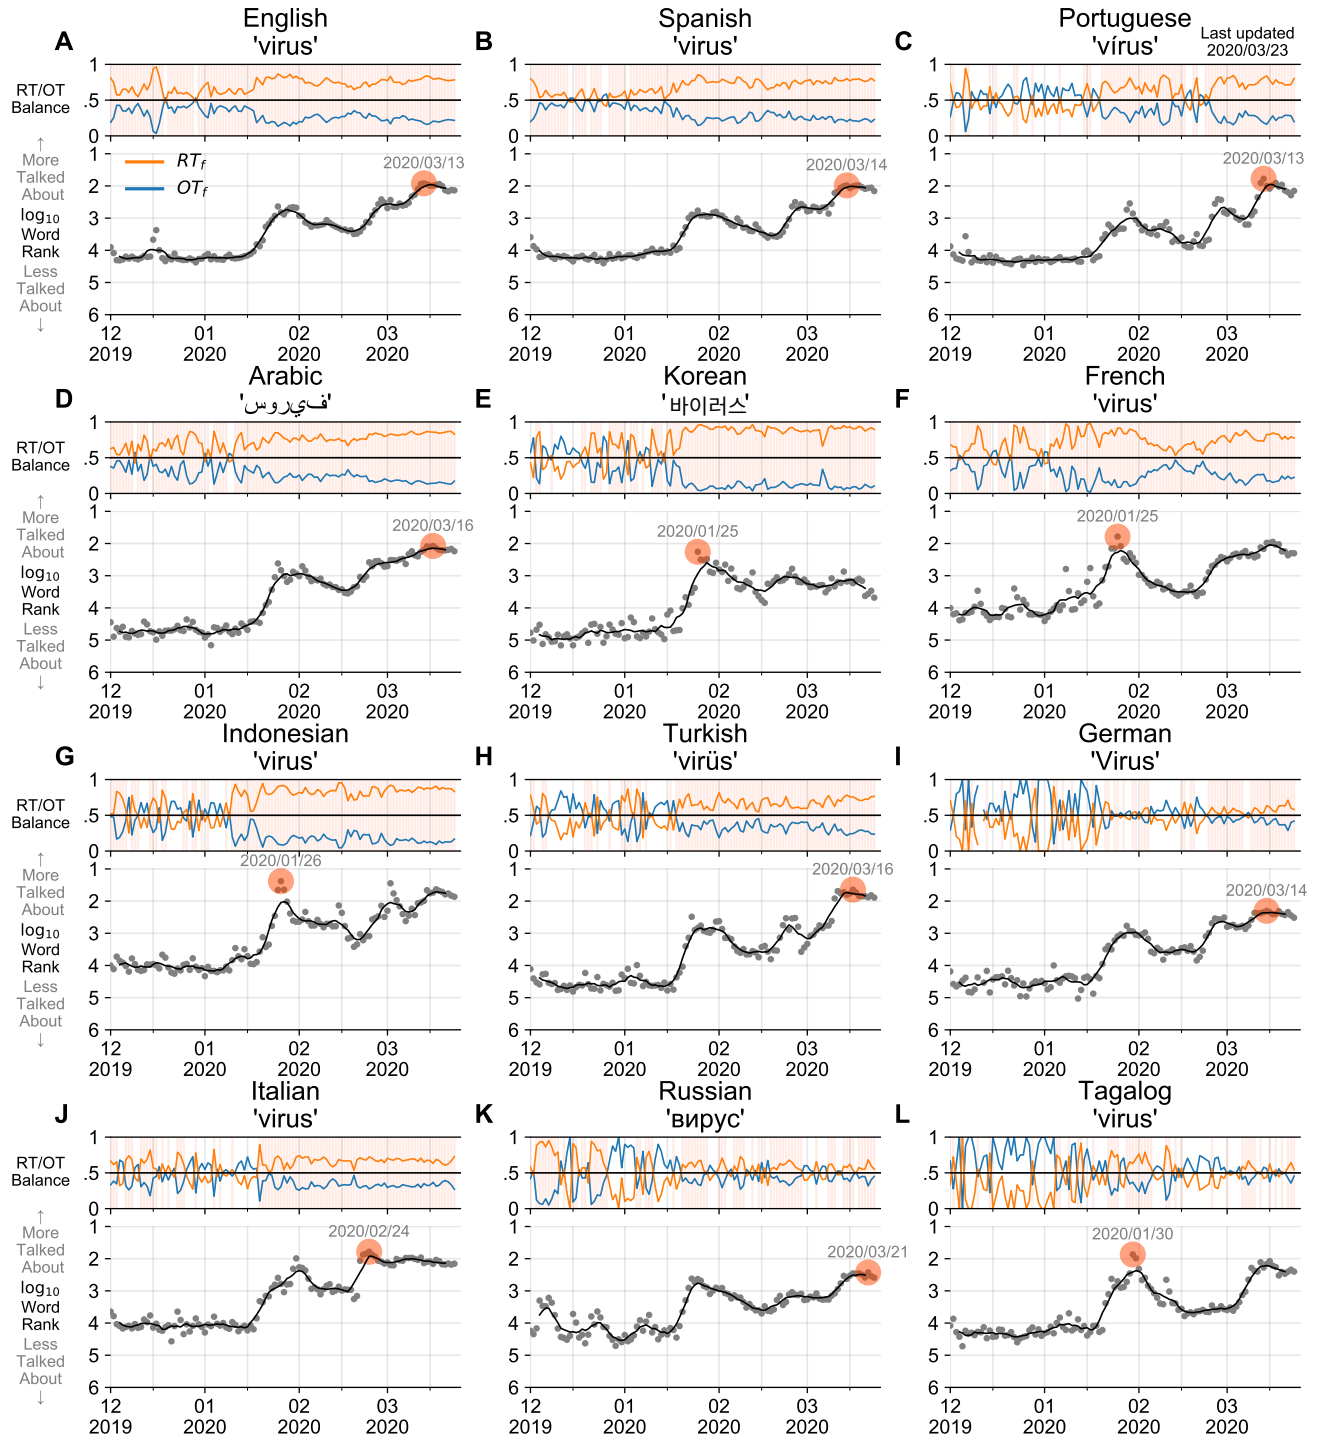


FIG. 2. Usage rank time series for the word 'virus' in the top 12 of the 24 languages we study here. **The major observation is that the world's attention peaked early in late January around the news of an outbreak of a new infectious disease in Wuhan, declining through well into February before waking back up.** The main plots in each panel show usage ranks at the day scale (GMT) (we use common sense ranking meaning the most frequently used 1-gram has rank $r = 1$). The solid lines indicating smoothing with a one week average (centered). The plots along the top of each panel show the relative fractions of each 1-gram's daily counts indicating as to whether they appear in retweets (RT, spreading) or organic tweets (OT, new material). The background shading shows when the balance favors spreading—story contagion. See Fig. 3 for the next 12 languages, as well as Sec. III for general discussion. See also the companion plots of relative frequency time series in Fig. A1.

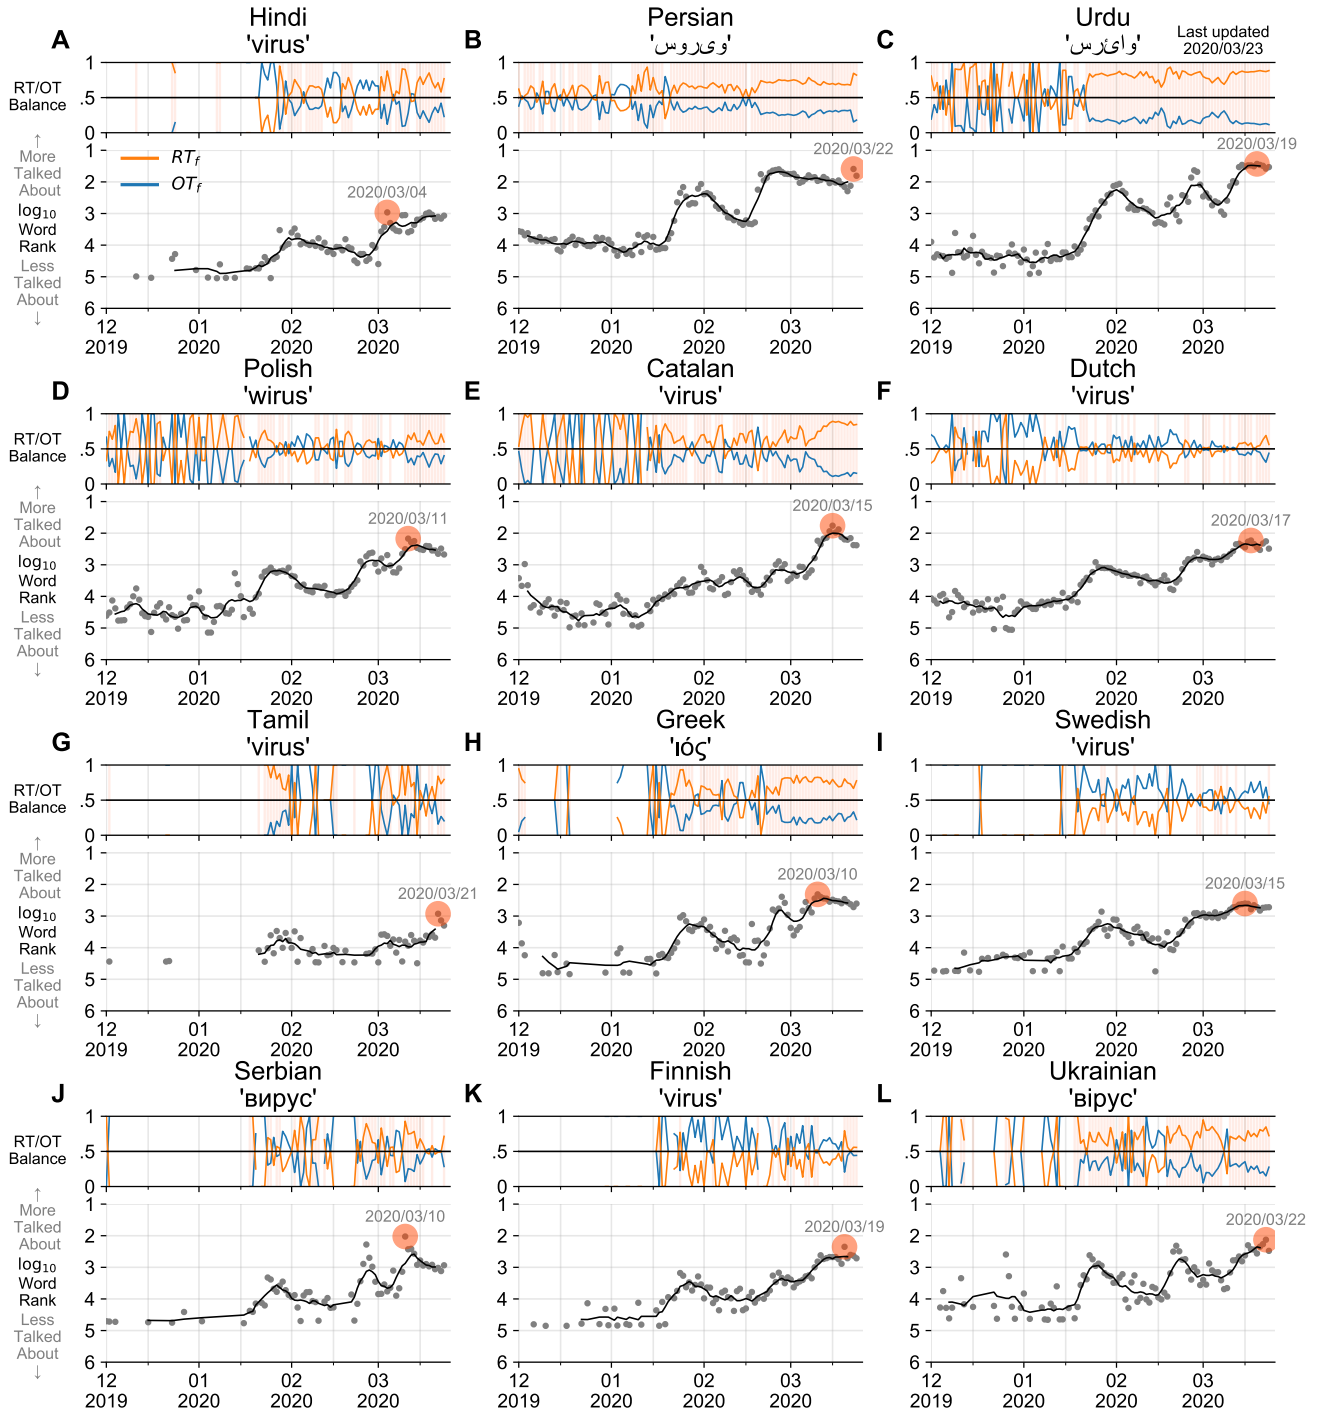


FIG. 3. Following on from Fig. 2, usage rank time series for the word 'virus' in the second 12 of the 24 languages. See also the companion plots of relative frequency time series in Fig. A2.

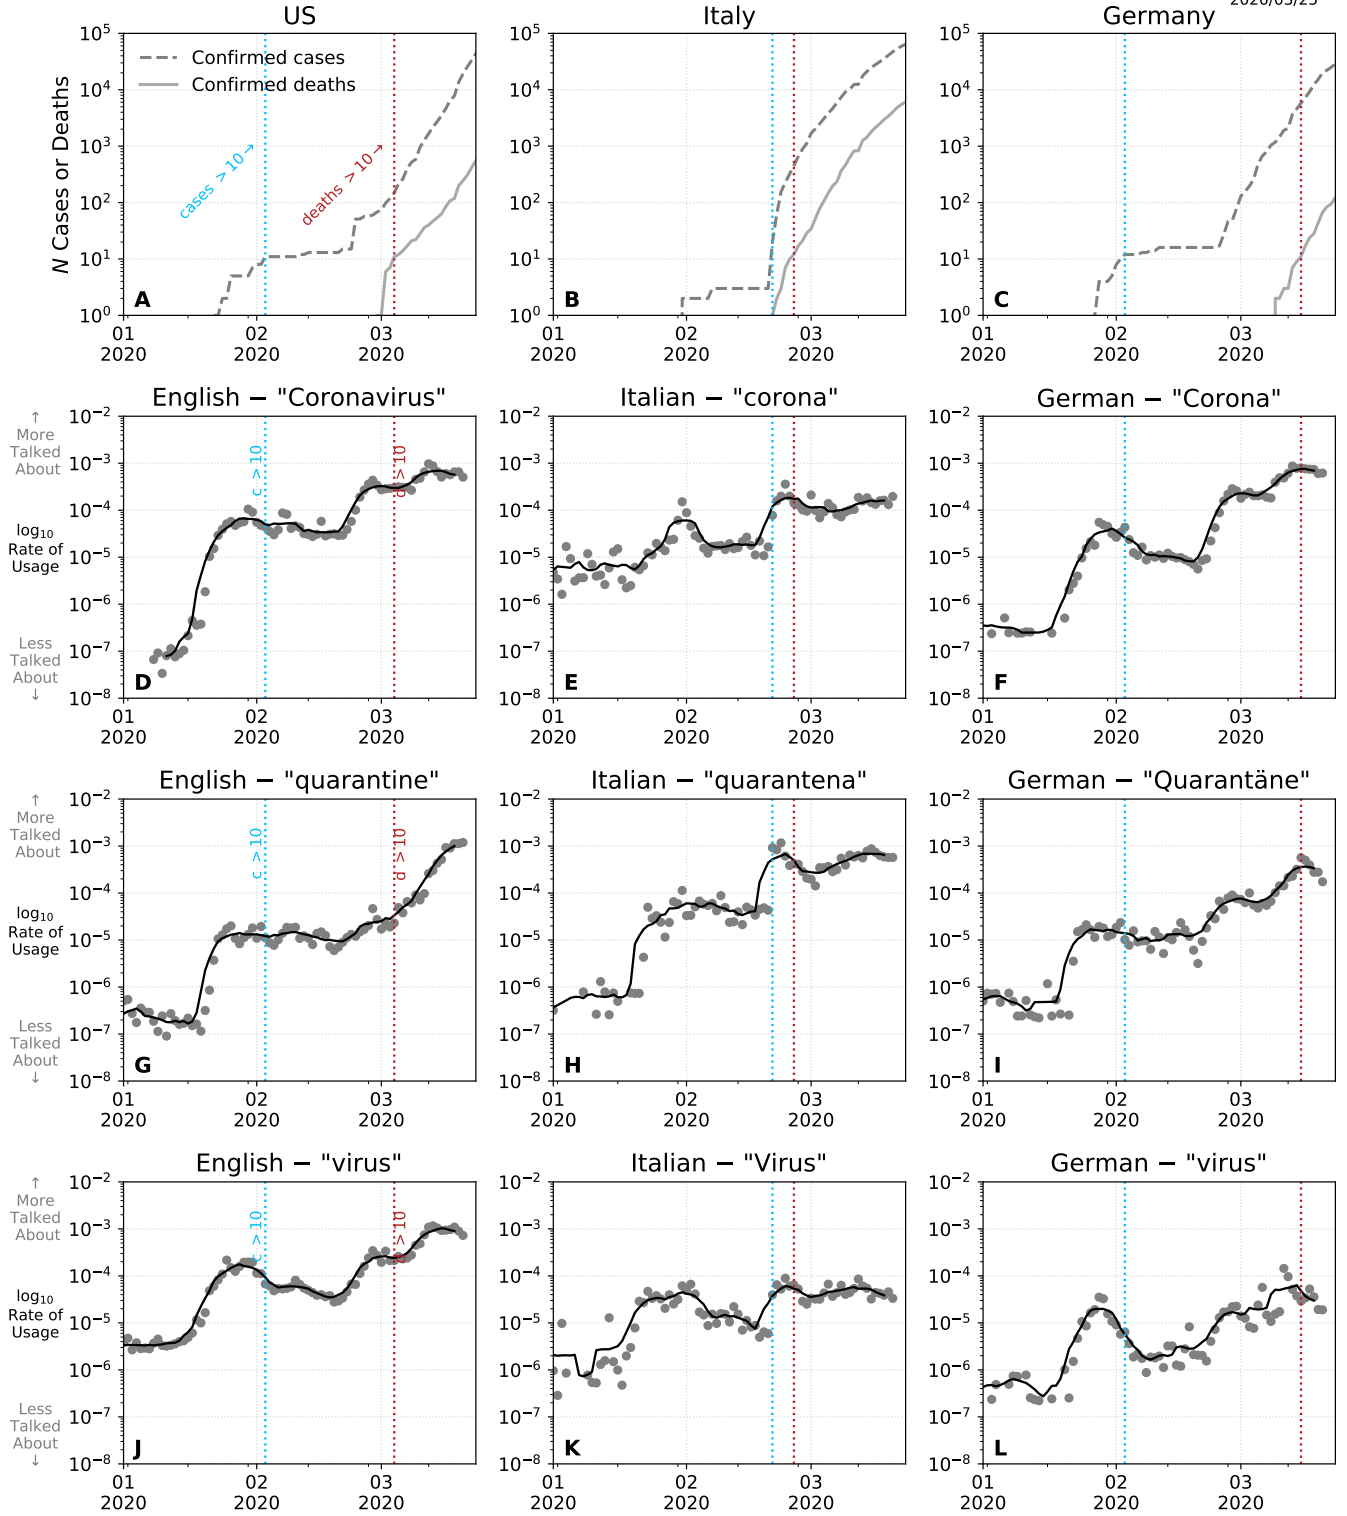


FIG. 4. Time series for known case loads and death counts for the United States, Italy, and Germany, compared with 1-grams for ‘Coronavirus’, ‘quarantine’, and ‘virus’ in English, Italian, and German. We note that the counts in plots A, B, and C are underestimates, more so for cases than deaths, and errors are unknown. We sourced data for confirmed cases and fatalities from Johns Hopkins University Center for Systems Science and Engineering’s COVID-19 project [3]. Starting on 2020/01/22, the project’s data has been collected from national and regional health authorities across the world. The data is augmented by case reports from medical associations and social media posts—these later sources are validated against official records before publication. For the present piece, we use daily summary files for case counts and fatalities, although an API and online dashboard are available for more up-to-date reports.

where testing has also been limited [3] and with peak attention on Twitter coming in January and early focus on economic issues and evacuation of nationals from Wuhan, a dip in the rank of ‘virus’ in the second half of February is also worrying (Fig 2G).

As one very simple example of comparing our Twitter times series with pandemic-related data, in Fig. 4, we present plots of confirmed cases and deaths over time for the United States, Italy, and Germany, along with time series for ‘Coronavirus’, ‘quarantine’, and ‘virus’ (translated for Italian and German). We indicate where numbers of known cases and known deaths reach 10 with vertical dotted lines, and note that these are surely under-reportings of unknown error. The points which known case loads and known death tolls reach 10 vary considerably across the three countries, with Italy rapidly moving from 10 known cases to 10 known deaths.

IV. DATA AND SITES

We share and maintain all data on Gitlab at: <https://gitlab.com/compstorylab/covid19ngrams>.

For each of the 24 languages, we provide time series for the top 1000 1-grams for both all tweets (AT) and organic tweets (OT). Note that in general, a 1-gram in the AT group may not appear in the OT group, as we start with only 10% of all tweets—a singular, highly retweeted tweet may contain 1-grams that are not produced organically by other users on that day (e.g., a misspelling by a famous user).

We include some further figures as supplementary material.

We also provide a connected website associated with our paper at:

<http://compstorylab.org/covid19ngrams/>.

We show tables of the leading 1-grams in our data set, as well as an example bar chart race for the dominant COVID-19 1-grams in English.

Our intention is to automatically update the data set on Gitlab, as soon as we have processed all tweets for a day.

V. CONCLUDING REMARKS

We echo our main general observation of how COVID-19 has been discussed through late March 2020: **After reacting strongly in late January to the news that a coronavirus-based disease was spreading**

in China, attention across all but 2 of the 24 languages we survey dropped through February before resurging in late February and through March.

We see abrupt shocks in time series as populations shifted rapidly to heightened levels of awareness, particularly in the Italian time series. In the time series for ‘virus’, we see two and sometimes three peaks of attention in the space of just a few months.

Our hope is that our collection of Twitter 1-gram time series that are especially relevant to March 2020 will be of benefit to other researchers.

The time series we share will, in part, reflect many other aspects beyond mentions of ‘virus’, which we have only briefly explored here. **Possible topics to investigate include washing (including the soap emoji), testing, serology, vaccine, masks and protection equipment, social and physical distancing, terms of community support versus loneliness and isolation, closures of schools and universities, economic problems, job loss, and food concerns.**

We repeat that the lists we provide are meant to represent the important 1-grams of March 2020, and we urge a degree of caution in the use of the data set.

As we have indicated above, our lists of 1-grams contain some peculiarities that will not be directly relevant to COVID-19. Entertainment (e.g., movies, celebrities, and K-pop) and sports (football along with sports in the United States) are standard fare on Twitter when no major events are taking place in the world. The extent to which these aspects of Twitter are submerged as pandemic related 1-grams rise is of interest.

Finally, while we have been able to identify languages well, geolocation is coarse and at best will be at the level of countries. The strength of geolocation for our time series will depend on the degree of localization of a given language as well as Twitter user demographics. We leave producing n -grams with serviceable physical location as a separate project.

ACKNOWLEDGMENTS

The authors are grateful for support furnished by MassMutual and Google, and the computational facilities provided by the Vermont Advanced Computing Core. The authors appreciate discussions and correspondence with Aaron Schwartz, Todd DeLuca, Nina Safavi, and Nicholas Danforth.

-
- [1] R. Li, S. Pei, B. Chen, Y. Song, T. Zhang, W. Yang, and J. Shaman, *Science* (2020).
 - [2] M. U. G. Kraemer, C.-H. Yang, B. Gutierrez, C.-H. Wu, B. Klein, D. M. Pigott, L. du Plessis, N. R. Faria, R. Li,

W. P. Hanage, J. S. Brownstein, M. Layan, A. Vespignani, H. Tian, C. Dye, O. G. Pybus, and S. V. Scarpino, *Science* (2020).

- [3] E. Dong, H. Du, and L. Gardner, The Lancet Infectious Diseases (2020).
- [4] B. Xu, M. U. Kraemer, B. Gutierrez, S. Mekaru, K. Sewalk, A. Loskill, L. Wang, E. Cohn, S. Hill, A. Zarebski, *et al.*, The Lancet Infectious Diseases (2020).
- [5] M. Cinelli, W. Quattrocioni, A. Galeazzi, C. M. Valensise, E. Brugnoli, A. L. Schmidt, P. Zola, F. Zollo, and A. Scala, “The covid-19 social media infodemic,” (2020), available online at <http://arxiv.org/abs/2003.05004>.
- [6] E. Chen, K. Lerman, and E. Ferrara, arXiv preprint arXiv:2003.07372 (2020).
- [7] V. Lampos, S. Moura, E. Yom-Tov, I. J. Cox, R. McKendry, and M. Edelstein, “Tracking COVID-19 using online search,” (2020), available online at <http://arxiv.org/abs/2003.08086>.
- [8] T. Alshaabi, D. R. Dewhurst, J. R. Minot, M. V. Arnold, J. L. Adams, C. M. Danforth, and P. S. Dodds, “The growing echo chamber of social media: Measuring temporal and social contagion dynamics for over 150 languages on twitter for 2009–2020,” (2020), available online at <http://arxiv.org/abs/2003.03667>.
- [9] P. S. Dodds, J. R. Minot, M. V. Arnold, T. Alshaabi, J. L. Adams, D. R. Dewhurst, T. J. Gray, M. R. Frank, A. J. Reagan, and C. M. Danforth, “Allotaxonomy and rank-turbulence divergence: A universal instrument for comparing complex systems,” (2020), available online at <https://arxiv.org/abs/2002.09770>.
- [10] A. Joulin, E. Grave, P. Bojanowski, and T. Mikolov, “Bag of tricks for efficient text classification,” (2016), available online at <http://arxiv.org/abs/1607.01759>.
- [11] A. Joulin, E. Grave, P. Bojanowski, M. Douze, H. Jégou, and T. Mikolov, “Fasttext.zip: Compressing text classification models,” (2016), available online at <http://arxiv.org/abs/1612.03651>.
- [12] E. A. Pechenick, C. M. Danforth, and P. S. Dodds, Journal of Computational Science **21**, 24 (2017), available online at <http://arxiv.org/abs/1503.03512>.
- [13] P. S. Dodds, J. R. Minot, M. V. Arnold, T. Alshaabi, J. L. Adams, D. R. Dewhurst, A. J. Reagan, and C. M. Danforth, “Fame and Ultrafame: Measuring and comparing daily levels of ‘being talked about’ for United States’ presidents, their rivals, God, countries, and K-pop,” (2019), available online at <https://arxiv.org/abs/1910.00149>.

APPENDIX

Language	# 1-grams
English (en)	3.75×10^9
Spanish (es)	1.41×10^9
Portuguese (pt)	1.22×10^9
Arabic (ar)	8.34×10^8
Korean (ko)	1.10×10^9
French (fr)	2.18×10^8
Indonesian (id)	3.32×10^8
Turkish (tr)	2.06×10^8
German (de)	1.17×10^8
Italian (it)	1.15×10^8
Russian (ru)	2.73×10^7
Tagalog (tl)	1.23×10^8
Hindi (hi)	8.05×10^8
Persian (fa)	1.22×10^8
Urdu (ur)	1.99×10^8
Polish (pl)	7.39×10^7
Catalan (ca)	1.37×10^8
Dutch (nl)	1.59×10^8
Tamil (ta)	1.30×10^8
Greek (el)	2.81×10^7
Swedish (sv)	1.91×10^7
Serbian (sr)	2.79×10^7
Finnish (fi)	1.34×10^7
Ukrainian (uk)	2.22×10^7

TABLE A1. Approximate number of occurrences of all 1-grams captured in our dataset by language for 2019/09/01 through to 2020/03/23. Languages are sorted by overall language popularity on Twitter for 2019.

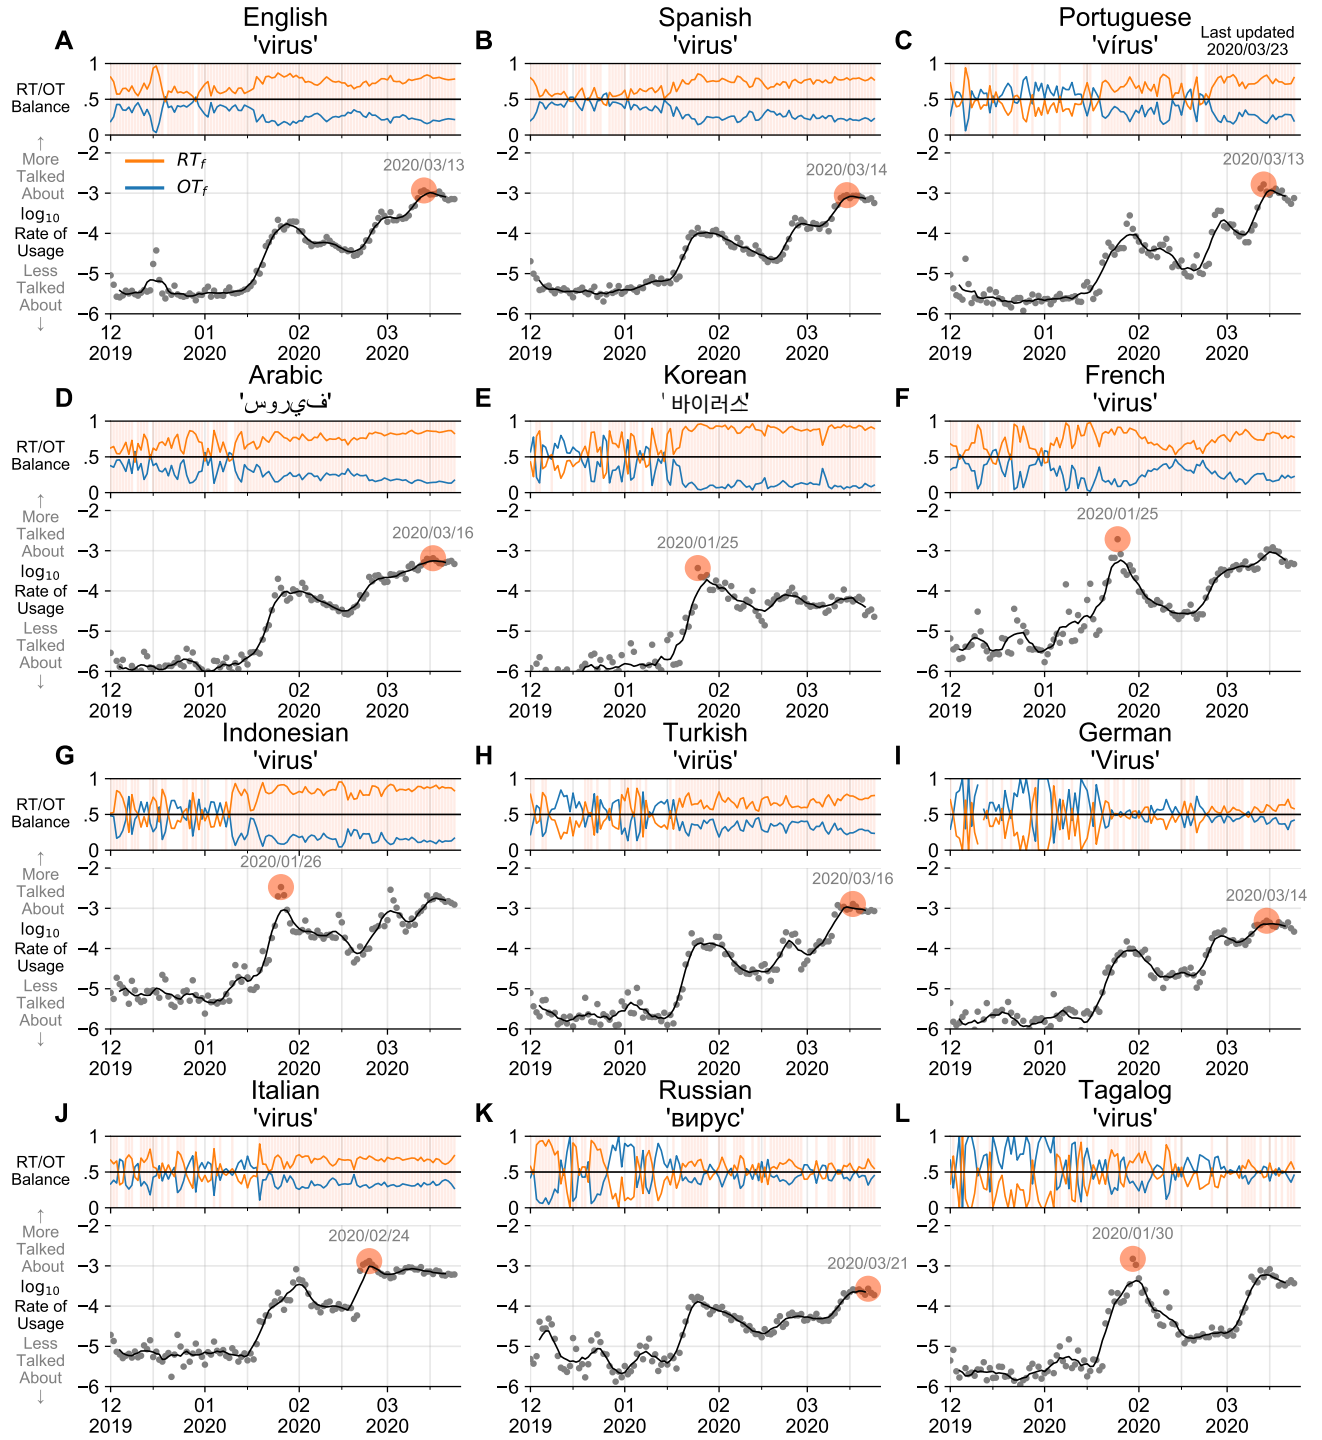


FIG. A1. Relative usage rate time series for the word 'virus' in the top 12 of the 24 languages we consider here, with language ranking by total usage in 2019. See the caption of Fig. A1 as well as Sec. III for explanations of the plot formats. See the companion plots in Fig. 2 for Zipfian rank time series.

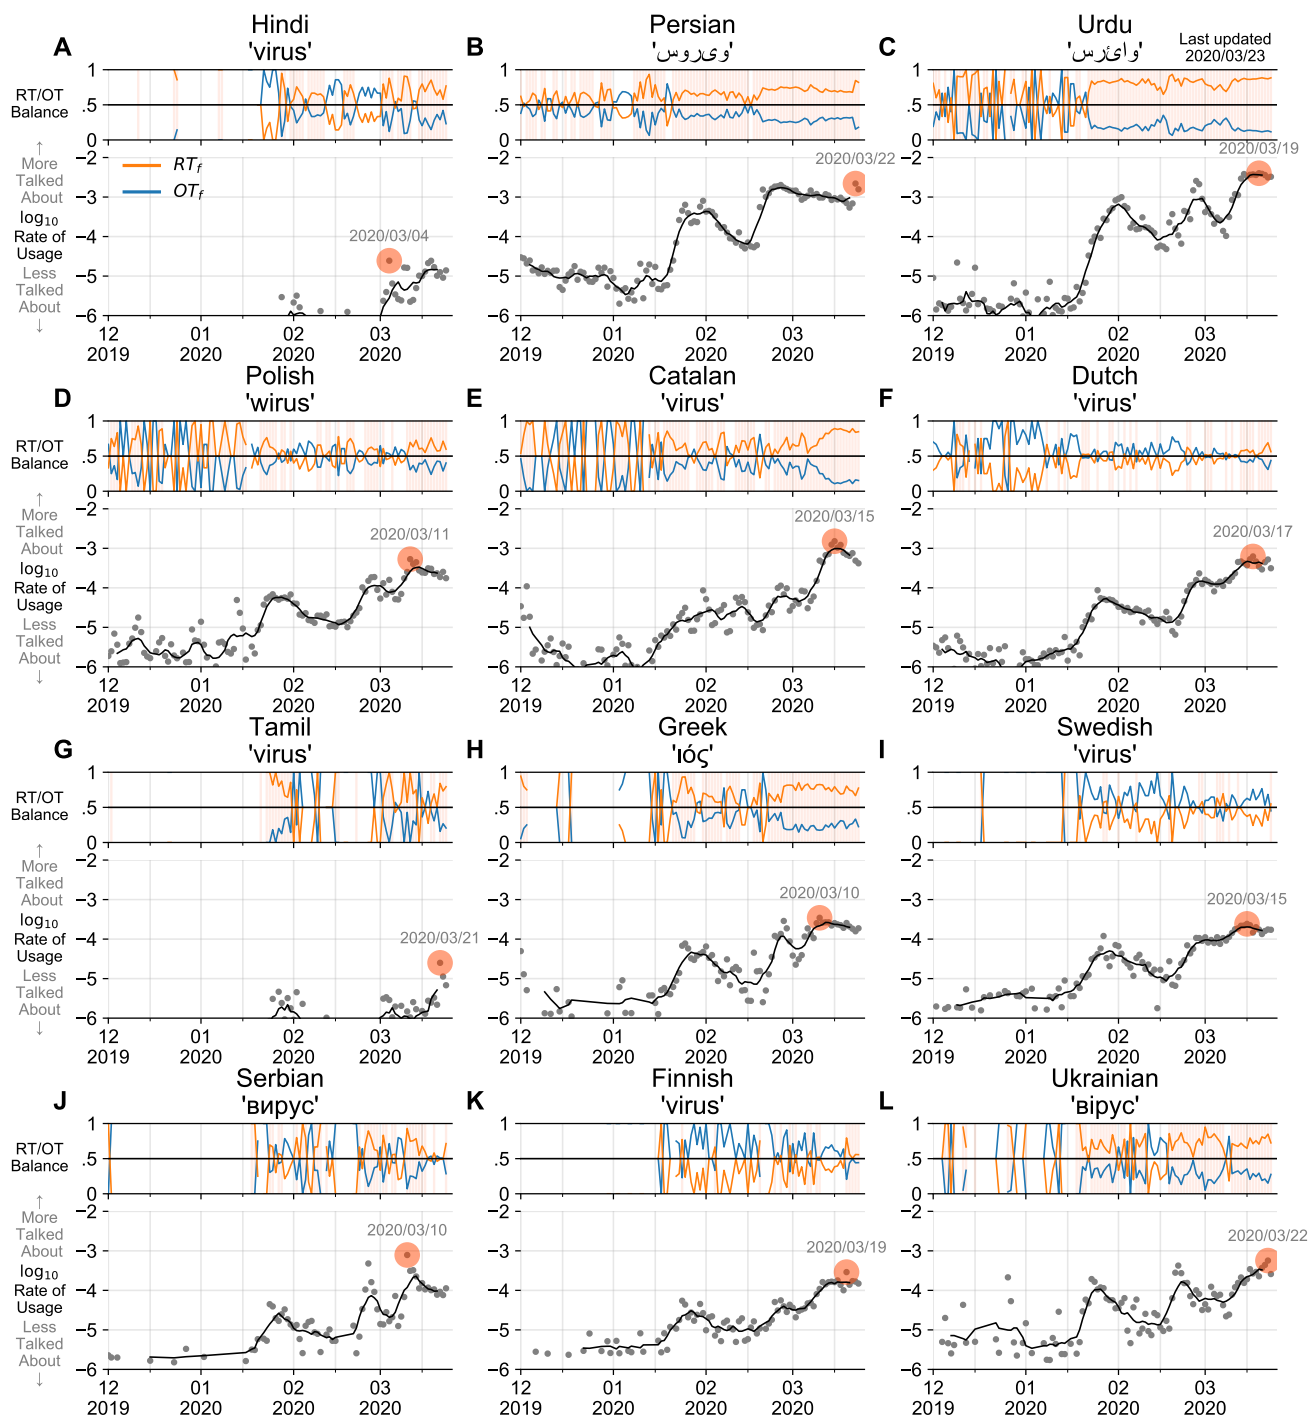


FIG. A2. Continuing on from Fig. A1: Relative usage rate time series for the word 'virus' in the bottom 12 of the 24 languages in our data set. See the companion plots in Fig. 3 for Zipfian rank time series.